

Evaluating composite events using shared trees

D. Moreto and M. Endler

Abstract: Distributed and concurrent systems tend to generate large numbers of events which require further processing, such as filtering and establishment of co-relations, to become useful for human or automated managers. The work focuses on the detection of composite events defined through event expressions involving primitive events and some operators: Boolean conjunction and disjunction; sequence, repetition, and absolute, periodic and relative timer events. A concrete result of this work was the design and implementation of a general purpose event processing service (EPS) which may be used by monitoring applications to be informed about the occurrence of primitive and composite events. Composite events are commonly represented as trees where the leaf nodes represent primitive event types and the intermediate nodes represent any of the supported event operators. The main contribution of this work is the description of a method to process composite events that share common sub-expressions with other composite events. The EPS was implemented in Java within the framework of the Sampa project.

1 Introduction

Collection and analysis of events are intrinsic tasks of monitoring computer systems and may be used for a wide variety of purposes, such as management of resource usage, detection of program state transitions, capturing of user interactions, performance analysis, security, etc. Thus, event processing has been used in many areas such as active databases, computer networks, software development tools, distributed systems, and many others.

In networked and distributed systems, monitoring has always been a very important tool for the management of resource utilisation and sharing, as well as for the verification of proper functioning of a system [1]. But monitoring also plays an important role in testing, debugging and optimising concurrent (i.e. parallel and distributed) programs, where a developer has to deal with many sorts of indeterminism. For most applications, monitoring also includes some form of event processing, such as filtering and detection of co-relations between events.

Composite events are essential for applications where it is impossible to detect a particular state of a system by looking only at isolated events. For example, in local area network management, a problem with a physical link may be detected when the amount of incoming IP packets *and* the amount of outgoing IP packets at the attached network interfaces ceases to grow (e.g. no IP/ICMP activity). Although the freezing of the number of packets processed at one interface is an isolated event that may or not be relevant information, in this particular case it is the combination of *several* such events that is of relevance for link-layer network management.

Similarly, monitoring and correlating different types of events at different hosts, such as, for example, large numbers of open TCP connections to/from a given set of hosts, or unusual processes with *root* permission may also be important for detecting abnormal system or user behaviours.

Another class of applications that rely on the detection of composite events are process control systems, where the environment is continuously monitored through several sensors (e.g. for temperature, pressure, motion, etc.). In these systems primitive events are generated by those sensors whenever some of its input variable trespasses certain boundary values, and many control (re)actions depend on the detection of patterns of events that have some application-specific correlation, such as

```
if (temperature(X or pressure)Y)
    and tank_is_full then start_reaction
```

A concrete result of this work was the design and implementation of a general-purpose event processing service (EPS) which may be used by applications which require some form of monitoring. Through EPS these can be informed of the occurrence of a composite event, i.e. a set of related events occurring at different elements of a distributed system. Monitoring applications submit requests for notification (*R/N*) to the EPS which contain an event expression and a condition on event attributes that must be satisfied for the application to be notified. The EPS also receives messages (event instances) from sensors notifying it of the occurrence of changes in either software or hardware components in the system. A sensor may be any system component capable of sending primitive event instances to the EPS.

Composite events are defined through an *event expression* based on primitive event types and combined by any of the following operators: boolean conjunction and disjunction; sequence and repetition, with arbitrary nesting degree. Moreover, it is possible to specify also periodic and absolute timer events and to relate the former to other primitive or composite events. Composite events are commonly represented as trees where the leaf nodes

© IEE, 2001

IEE Proceedings online no. 20010241

DOI: 10.1049/ip-sen:20010241

Paper first received 29th November 1999 and in revised form 5th January 2001

The authors are with the Departamento da Ciência da Computação, IME-University of São Paulo, Rua do Matão 1010, 05508-900 São Paulo, Brazil.

E-mail: {moreto,endler}@ime.usp.br

IEE Proc.-Softw., Vol. 148, No. 1, February 2001

represent primitive event types and the intermediate nodes represent any of the supported event operators.

The main contribution of this work is the development of a method to process composite events that share common subexpressions with other composite events. Unlike other event processing systems [2, 3] which do not support aggregated trees, this method allows for an optimised evaluation of event expressions since less leaf nodes need to be checked when a new event instance arrives from a sensor.

To support the sharing of subtrees (derived from common subexpressions within different $RfNs$), our method also takes into account that $RfNs$ may specify different conditions on the attributes of event types appearing in the common subexpressions. This makes the accounting of event usage for the different conditions more complicated, since for example, a primitive event instance that has not contributed to the satisfaction of one condition cannot be eliminated unless there is no other event condition to which it may contribute.

Our method consists of a four-phase tree-traversal algorithm in which the first phase is executed whenever a new event instance arrives at the EPS, the second and the third phases occur whenever an event condition of an RfN is satisfied, and the fourth phase can be done at any moment and aims at garbage-collecting the events that actually contributed to the satisfaction of a condition. The algorithm used in our method has been formally described in an attribute grammar and has been implemented in the event processing service. Despite being a general-purpose service, EPS is best suited for applications that require the detection of complex event patterns (eventually with similar structures), such as found in workflow control, network management or monitoring of automated processes (e.g. a chemical reaction, manufacturing, etc.).

This work was done in the framework of a larger project aiming at the development of a system for supporting the execution and management of fault-tolerant distributed applications and services. This project is called Sampa (*system for availability management of process-based applications*) [4], and also comprises other component services, such as monitoring sensors, checkpointing, group communication and configuration management, which are used to monitor, control and dynamically configure systems in response to the detection of failures. In this project, event monitoring is a basic service which is to be used for the purpose of availability management. EPS is also being used as a basic component for exception handling in an object-oriented software architecture for dependable systems [5], within another project.

2 Related work

Most of the works dealing with composite events are related to the area of active database and are based on centralised monitors. Only few works, such as GEM and EVE, are for distributed systems. While most of the works use trees as the basic data structure for implementing its event monitors, other computational models are used as well, such as Petri nets or finite automata.

Samos [6] is a project that uses a colored Petri net model to specify the behaviour in active databases. This model is very suitable for modelling and describing complex and concurrent behaviour, but its implementation was shown to be complex and comparatively inefficient [7].

ODE [8] is a project in object-oriented databases that was developed at AT&T Bell Labs. Its main characteristic is the use of triggers to automatically launch actions depending on the state of the database. One of the results of the project is COMPOSE, a system for specification and detection of composite events which is based on finite automata. Its main weakness is the impossibility to detect concurrent events and to reorder events according to their creation times.

The following projects use trees to represent composite events: EVE [2] is an event detector for workflow control; Snoop [9] is a monitor for databases and GEM [3] is a generic event monitor for distributed systems. The main reason for adopting trees is that they are simple, well suited to represent composition, and have been thoroughly studied in the past, leading to a large collection of efficient algorithms for traversal and manipulation.

In the monitors of EVE and GEM each request-for-notification (RfN) for a composite event causes the creation of an entire new event tree, even when this tree contains identical subtrees of other trees already created by the monitor. Consequently, the monitor may process a same subtree several times, depending on the number of $RfNs$ for similar composite events. For example, in a local area network it may happen that several programs, such as a performance or security management tool, a distributed operating system or a middleware service may be interested in notifications about very similar composite events concerning, for example, TCP connections (or IP traffic) to or from a specific host.

Snoop is a monitor which is also concerned with the problem of repetitive processing of shared subtrees of composite events, and also uses aggregate trees to implement composite event detection. Unlike our approach, however, Snoop defines a predefined set of different modes (called *contexts*) for processing composite events. The set consists of the following four processing modes, which will be explained considering a generic composite event type $E_C = (E_1, E_2)$ with constituent event types E_1 and E_2 .

Recent: only the most recent instances of the events are considered during the processing, i.e. when several instances of a given event type (e.g. E_1) have been generated, then all but the most recent instance are ignored.

Chronological: the instances are processed according to the order in which they have been generated. Thus an instance of a composite event type E_C is always built from the pair of the oldest instances of the constituent event types E_1 and E_2 . After the composite instance is processed, its constituent instances are not used again, and the next oldest pair of instances of E_1 and E_2 is considered, and so on.

Continuous: in this mode the instance of the constituent event type (e.g. E_1) that triggered the detection of the composite event E_C is combined with all the instances of E_2 . All the other instances of E_2 used in the composition are removed and will not be used in future instances of E_C . The instance of E_1 that triggered the detection of E_C may or not be used in future instances of E_C , which depends on the composition operator. For example, for operator conjunction it is reused while for the operator sequence it is removed.

Cummulative: when a composite event is detected in this mode, the instance of E_C contains all the instances of both E_1 and E_2 that have been generated so far, and the instance sets for this event types are reset.

3 System model

This work assumes a distributed application as being composed of processes executing on different hosts interconnected through a reliable communication network.

Some of the processes (the *sensors*) are exclusively dedicated to probe run-time status information about other processes or system resources. We also assume the existence of one (or several) monitoring components (the *monitors*), whose task is to do the application-specific analysis and presentation of the information provided by the sensors. This information is transferred via *event messages* (of a certain type) with a time-stamp.

In this model we further assume that clocks are approximately synchronised (within a certain precision). Due to this assumption, and because sensors do not necessarily interact with each other, we chose to use time-stamps with real-time values instead of vector time-stamps to determine the precedence of event messages at arrival. Thus in EPS time-stamps are supposed to be the value of the host's local clock at the moment the message is sent by the sensor.

4 Event processing service

The event processing service is a general-purpose service that supports the detection of primitive and composite events for any application which requests this service. An application process makes a request by sending a *RfN* to the EPS. The *RfN* is composed of the following parts:

Event expression: It specifies the structure of the composite event in terms of some primitive event types and five possible event operators, described in Section 5. The set of all pre-defined primitive event types is a parameter of the EPS, and each event type associated with a set of specific attributes.

Condition expression: It determines constraints on attributes of event types occurring in the expression which have to be satisfied for a notification to be sent to the requesting application process.

Callback port: This is a reference to an entry point (e.g. a port or remote method) at the requesting application process, which is accessed/invoked by EPS to send a notification of occurrence of an event of interest.

Notification: This is a character string that the requesting application process will receive from EPS as part of the event notification. It will typically describe the application-specific meaning of the corresponding event occurrence.

Each request for notification has a unique *request-Id*, which may further be used by the application process to access it, as for example, to cancel a particular request. The monitoring of the system is actually performed by application specific *sensors* (i.e. software processes), which report the occurrence of any significant change in either software or hardware component, by sending event messages to EPS. Each such message contains a single primitive event instance with some attribute values.

4.1 Architecture

The EPS is implemented as a centralised server program (shown in Fig. 1) composed of the following four functional units:

Tree management: This unit handles the acceptance of new *RfNs*, adds the corresponding event nodes in the event tree and sends requests for timer events to the timer unit.

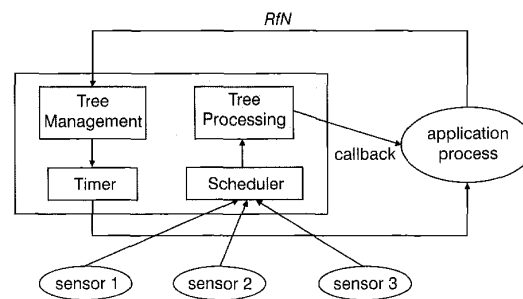


Fig. 1 Structure of event processing service

Scheduler: The scheduler receives event messages (with primitive event instances) from the sensors, implements a detection window to eliminate outdated events (Section 7.1) and schedules the delivery of event instances to the **tree processing** unit. To deliver the event instances in the correct order, the scheduler uses the *scheduler time* (cf. Section 7) to wait for “late” instances.

Tree processing: This is the core unit within EPS. It performs the entire event processing, as well as garbage collection on one or more trees representing the event expressions specified in the *RfN*. At every delivery of a new primitive event instance by the **scheduler**, this unit eventually generates new event instances at intermediate nodes of the tree to represent the occurrence of composite events. The algorithm for tree traversal used in this unit is described in detail in Section 6.

Timer: This unit essentially implements an alarm clock, which receives requests to generate either absolute or periodic timer events. These requests arrive either from the **tree management** or the **tree processing** units, where the first one requests absolute and periodic timer events and the second one requests absolute timer events.

5 Event types and attributes

Events are instances of either *primitive* or *composite event type*. Primitive event types have attributes that hold data values reflecting some change of state or behaviour in either a software or hardware component of the distributed system. Composite event types describe complex events, such as the conjunction, disjunction, sequence, etc. of simpler primitive or composite event types. A composite event is defined by an event expression contained in a *RfN*.

Common to all event types are the attributes **origin** and **time-stamp**. While for primitive events the value of a time-stamp is determined by the generating sensor, for composite events the time-stamp is inferred from the time-stamps of its component events.

To combine primitive events EPS has five operators: operator “&” (boolean and); operator “||” (boolean or) operator “;” (**sequence** of events); operator **TimesEq** and operator **TimesDf**. The last two operators represent the repetitive occurrence of a specific event. While operator **TimesEq** represents the repetition of an event type whose instances have the same value for a specific attribute, operator **TimesDf** does the opposite, it considers only the instances with different values for a specific attribute.

In addition to these operators it is possible to group an event expression by using parenthesis. For example, the event expression

$$TimesEq(n, E_1, attr_1) || (E_2; E_1)$$

describes the composite event of n occurrences of a event E_1 (all with the same value of attribute $attr_1$) or the composite event “ E_2 followed by E_1 ”. Such an event expression could be part of a RfN .

Although the ability to express the absence of events certainly adds expressive power to an event description language, we chose not to support negative event expressions in EPS, since this would require the implementation of complex evaluation methods based on interval logic [10].

Moreover, one can get around this problem by using special timer-equipped sensors (or interposing them between the normal sensors and EPS) that periodically check if some state change or event has occurred within some period of time, and if this is not the case, deliver an event to EPS of the type “not event E within last T time units”. Thus we believe that for most applications the lack of negative event expressions will not be a major problem.

In addition to the composite operators, EPS is able to deal with timer events. The three kind of timer events which EPS can handle are: **Absolute** (specifying a unique moment in time); **Periodic** (specifying a repetitive event) and **Relative**. This last event type defines the generation of a timer event some time after the occurrence of an arbitrary composite or primitive event. For example, $at(E, 0/0/2 0:01:05)$ specifies that exactly two days, one minute and five seconds after the occurrence of an event of type E a timer event will be generated.

6 Event processing

As in some other works we chose to represent a composite event as a tree, where the leaf nodes represent primitive event types (denoted by E_i) and the intermediate nodes represent one of the possible event operators presented in Section 5. In addition, there is a third kind of node called an *application node*, which represents an RfN and may be connected to a leaf or an intermediate node. Such event trees are built according to the event expressions of the $RfNs$ received by EPS from the application processes. Unlike most other work in this area, where each application request causes the creation of a separate tree representing the corresponding event expression, our approach aggregates the trees whenever an event expression from an RfN shares subexpressions with another RfN , eventually received from another application. The main advantage of this tree aggregation is that it facilitates the processing of events contained in several expressions. Only Snoop [9] provides similar functionality. However, unlike Snoop which adopts four modes to control the event processing, EPS uses a generic mode where all possible combinations of events are considered. Details about Snoop’s four modes and the advantage of using a generic mode in EPS are discussed in Section 10.

In what follows we make a distinction between *nodes* of the event tree, e.g. describing primitive event types and operators, such as operator “&”, operator “||”, etc. and *event instances* (or simply, *instances*), which are the pieces of monitored data obtained through event messages from the sensors. Whenever an instance arrives at the EPS, it is associated with the corresponding leaf node and its occurrence is also propagated up the tree. This upward propagation is done by eventually creating new instances at the parent nodes of the leaf node, at the parent nodes of these, etc. Thus while processing incoming event instances, several trees of instances are created according to the operators represented by each of the intermediate nodes of the event tree.

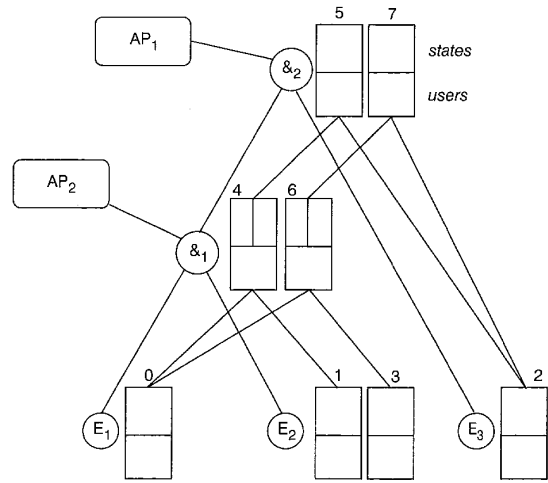


Fig. 2 First phase of event processing

Fig. 2 depicts a simple aggregated event tree built from two $RfNs$ sent by applications AP_2 and AP_1 . In this example the event expression from AP_1 is $(E_1 \& E_2) \& E_3$ and from AP_2 is $E_1 \& E_2$. Moreover, the Figure shows the instance trees corresponding to the arrival of four primitive event instances 0, 1, 2 and 3. In this and the following figures we represent the event trees by circles (the nodes) and thick lines, and the instance trees by rectangles (the instances) and thin lines. In the rectangles representing the instances, the upper and lower parts are used to represent the attributes **states** and **users** of each event instance, which are explained in the following.

Since we are using aggregate trees, each node can have one or more parent nodes. Therefore, for the event processing we need to decide if an instance has or not been already used in a composite event including a specific parent node. This decision is supported by inspecting the following two attributes associated with every instance e of a tree node E .

states_e is a set which contains one flag for each parent node of E . Each flag is either *marked* or *not_marked*, which indicates whether this instance has or has not been used in an application notification involving the corresponding parent node;

users_e is a set of identifiers that holds all ID applications which were notified of the occurrence of a composite event instance of which instance e is a component.

The event processing in EPS consists of four phases, each of which is responsible for the following actions: the first phase creates intermediate event instances; the second and third phases eventually notify the application and mark all instances which participated in the notification; and the fourth phase does garbage collection, by removing parts of (unusable) instance trees.

6.1 First phase

The first phase builds the instance trees. It is an incremental, bottom-up process, which happens every time that an instance is delivered to a leaf node by the scheduler. The arrival of such instance causes the creation of additional instances at the corresponding parent node. The number of new instances that are created at such a node depends on the particular operator represented by the node and on the

set of already existing instances at its child nodes. For example, at a node representing an operator “&” the instance set is the cartesian product of the instance sets of the corresponding child nodes. At an operator “||” node the instance set is the union of the instance sets at the child nodes. This is because any of these instances may be used as an instance of the parent node (representing operator “||”).

The event operator “;” is a particular case of the operator “&”, with the additional constraint that an instance from the right branch (right child node) must be combined only with instances of the left branch which have a lower time-stamp attribute. For example, if we assume that the left child of the node representing operator “;” has two instances, e_{11} and e_{12} with time-stamps $e_{11}.t=10$ and $e_{12}.t=12$ respectively, and a new instance, e_r with $e_r.t=11$, arrives at the node's right child, then, because of the precedence constraint, only one new composite instance, $e_{11};e_r$ is created at the composite node representing operator “;”.

At nodes representing either operator **TimesEq** or operator **TimesDf** the set of instances is the set of all possible choices of N elements out of the set L of instances at the (single) child node. Thus in this case we build an instance set with cardinality

$$\binom{card(L)}{N}$$

The above rules are used for the creation of new intermediate instances at all levels of the event tree except for the bottom level, which is composed of leaf nodes.

Notice that by our method we build composite instances by forming all possible combinations of the child instances. This generic form of composite event derivation (*All Combinations*) has some advantages compared to the more restrictive derivations defined by Snoop's [9] modes *recent*, *chronological*, *continuous* and *cumulative*, presented in Section 2.

Fig. 2 shows an event and the instance trees after processing the first phase, for instances 0, 1, 2 and 3.

6.2 Second phase

The second phase starts when a new set S of composite instances arrives at an application node representing an R/N . For each composite instance in S , the event condition of R/N is tested. If any composite instance satisfies the condition, the corresponding instance tree is traversed down, and all the instances which contributed to building the composite instance are marked. Otherwise, no action is taken.

For each new set S of composite instances that arrives at an application node, only one instance in S can trigger a notification to the application. This is because all composite instances in S have at least one primitive instance in common, but the application should be notified only once of the occurrence of this “common” primitive event. Accordingly, all other instances of set S (and its component instances) should be marked so that they are not used again, i.e. they are not able to cause new notifications. This marking is performed during the third phase (Section 6.3).

Fig. 3 shows how the instance trees look like after the second phase, assuming that instance 7 satisfied the event condition of an R/N , hence causing a notification to application AP_1 . The black dots in this Figure represent the marks of attribute **states** _{e} at the event instances, as explained.

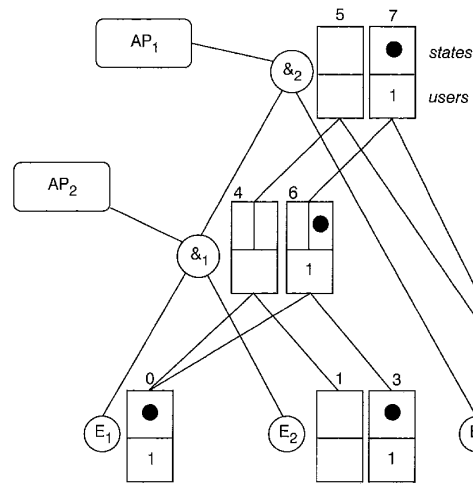


Fig. 3 Second phase of event processing

6.3 Third phase

In the second phase we could mark the instances which constituted the final composite event. However, many other instances use the same set or subset of the primitive instances of the final composite event. As mentioned in the Section 6.2, all of these instances are not supposed to generate new instances and therefore should be marked too. To mark all these instances, the third phase traverses the event tree up marking all instances that contain at least one primitive instance that is part of the composite event in S which caused the notification to the application (i.e. instance 7 in Fig. 3).

To ensure that no wrong instances are marked, the third phase checks only the instances of the nodes in the event tree corresponding to the specific event expression which caused the notification and initiated the second phase.

Fig. 4 shows the instance tree at the end of the third phase.

6.4 Fourth phase

The main goal of the fourth phase is to control the growth of the number of event instances. After a notification the instance tree may have many marked instances, and these




Fig. 4 Third phase of event processing

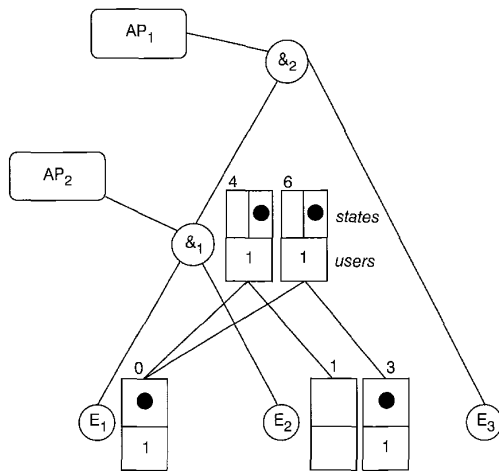


Fig. 5 Fourth phase (garbage collection)

instances cannot be used any more. The fourth phase traverses the event tree down, starting at the application nodes that have had a notification and removing all the instances which (i) have a mark in *all* elements of *states* and (ii) do not have any more a parent instance. Through this phase we guarantee that an instance is only kept in the tree if it can be part of a composite component for some future notification.

Fig. 5 shows the final instance tree after the whole event processing (phases one through four) caused by the occurrence of the primitive instances 0, 1, 2 and 3.

Although by this event processing method the first phase tends to generate large numbers of intermediate event instances, the marking and removal of instances performed in the phases two through four leaves only the minimum number of instances eligible for contributing to future notifications. This is very important for ensuring a good processing performance. Since EPS uses the generic processing mode (Section 6.1), the more instances are kept at a given level of the tree, the more composite instances at the upper levels will be generated from the levels below, yielding a direct impact on the monitor's performance.

7 Timing issues

Some well-known problems when monitoring distributed systems are the impossibility of achieving exact synchronisation among the clocks of the hosts in a network and the unpredictable delay of message transmission. These problems have been taken into account in the design of EPS, in particular, in the way it orders incoming event instances and handles communication delays.

7.1 Handling of delays

Concerning the first problem, in EPS we assume that machine clocks are approximately synchronised (with error rate $\leq \epsilon$) through the use of some time-synchronisation protocol [5]. Based on this assumption it makes sense to compare the timestamp of events (received from sensors) to determine an approximate relative precedence.

To deal with the problem of message transmission delay, EPS has the following two parameters:

detection window is a window size Σ (i.e. a period of time in *ms*) used to determine acceptance of event instances for

processing by EPS. An instance e can be accepted if its time-stamp attribute $e.t$ is greater or equal to $CT - \Sigma$, where CT is the current time at the arrival of the event instance. An appropriate choice for a value of Σ mainly depends on the kind of application which is using the EPS. A large Σ is recommended for applications with relatively long periods of stable states, i.e. where event instances are not generated frequently. Examples of such applications are monitors for chemical reactions. A small Σ is recommended for applications with frequent state changes, as for example, a monitor aiming to detect dependencies among concurrent computation tasks.

scheduler time is a delay τ in *ms*. This time is used to wait for event instances which eventually arrive at the EPS with some delay. This is particularly important in a distributed environment, where the event instances sent by the sensors may arrive in a different order than their generation times.

7.2 Event ageing

Another problem related to the monitoring of events is that some instances of primitive and composite events will never be marked by any application, mainly because their attributes do not satisfy the corresponding event conditions in the R/N .

To get rid of those events, every event processing system needs to implement a garbage collection based on event ageing. In the EPS, such garbage collection (based on event age) is done by traversing the tree top-down, scanning and removing all instances e which have their time-stamp attribute $e.t$ satisfying $e.t \leq CT - (1/n) * \Sigma$, where Σ is the detection window and n is the frequency in which the garbage collection is executed. During this traversal, whenever a composite instance is to be removed due to its age, all of its child instances can be removed as well. This is due to the fact that the time-stamp of a composite instance is always the maximum of the time-stamps of its child instances.

The frequency n of the garbage collection has significant influence on the performance of EPS. If the frequency is high, the number of event instances in the tree is lower, but each garbage collection will remove few instances. If the frequency is low, the instance tree will have more instances for a larger period of time, but each garbage collection will remove many instances. Thus a reasonable periodicity for age-based garbage collection must be found.

8 Comparison with related algorithms

We compare the algorithm for processing event trees used in EPS with those used in Snoop [9] and GEM [3] monitors. Unfortunately, no performance data of these other monitors have been published or made available, making it impossible to compare execution times for similar composite event expressions. Table 1 summarises the differences in the algorithms used in the monitors.

8.1 GEM

Both GEM and EPS use a single algorithm for processing the trees representing composite events. The main difference between GEM and EPS is that the former does not consider the sharing of subtrees. This may cause the construction of several copies of a tree representing the same composite event requested by different monitoring applications. Because of its simpler structuring of compo-

Table 1. Comparison between algorithms

Monitor	Trees	Instance ordering	Tree processing	Modes
EPS	shared	before tree processing	four phases	generic
Snoop	shared	no ordering	one phase	four modes
GEM	simple	during tree processing	two phases	generic

site events (i.e. without sharing), GEM has also a simpler tree processing algorithm which consists of two phases: in the ascending phase the composite instances are created, and in the descending phase used instances are removed. Essentially, these two phases correspond to EPS's first and fourth phases.

The main disadvantage of GEM's lack of subtree sharing is that the following two actions have to be carried out whenever a new primitive event instance is delivered to the monitor: first, all trees with the corresponding event type have to be identified and then the processing associated with that instance has to be carried out in all the corresponding trees.

To reduce processing at the different trees GEM employs *filters* that specify which instances are allowed to ascend the tree. This filter is a logical condition on the instance's attributes and may be associated with any node in the trees. By using these filters it is possible to reduce significantly the repetitive processing of instances in different trees.

Yet another important difference is related to the ordering of instances which arrive at the monitor. Unlike in EPS, where the ordering is performed before starting the first phase, in GEM the ordering is done during tree processing. This means that when an old instance arrives at GEM, it is put at the correct position with respect to the other instances at the same node.

8.2 Sentinel

As mentioned in Section 2 Snoop defines four modes for processing composite events. However, sometimes there exist composite events of interest to some applications which cannot be expressed by any of these modes, explained as follows.

Unlike Snoop, in EPS we chose to process events in a complete way, i.e. always building composite instances from all possible combinations of the instances at the child nodes. This generic processing mode is required when the instances are to be combined by considering only their *logical dependence* as expressed by the operators, rather than their temporal attributes. This form of generic processing is required by some applications where not only the most recent events are relevant for detecting and managing abnormal situations. Examples of applications are network security management or dataflow monitoring in a distributed system.

For example, consider the case in which the goal is to detect cycles in message transmissions among a set of processes, for example to avoid indefinite forwarding. In particular, consider the specific problem for three processes P_1 , P_2 and P_3 . Here we may have two types of events, *send type* and *receive type*, denoted by $[st, procID, msgID]$ and $[rt, procID, msgID]$, where *procID* and *msgID* are the attributes that indicate, respectively, the process at which the event occurred and a unique message identifier. In this case a cycle can be detected if a process receives a message that it has previously sent. This is essentially a sequence of the primitive event *send* followed by the primitive event *receive* with matching

attributes. Translating this into an event expression and condition, it means that we are interested in composite events of the form $[st, X, Y]; [rt, Z, W]$, where $X=Z$ and $Y=W$ hold.

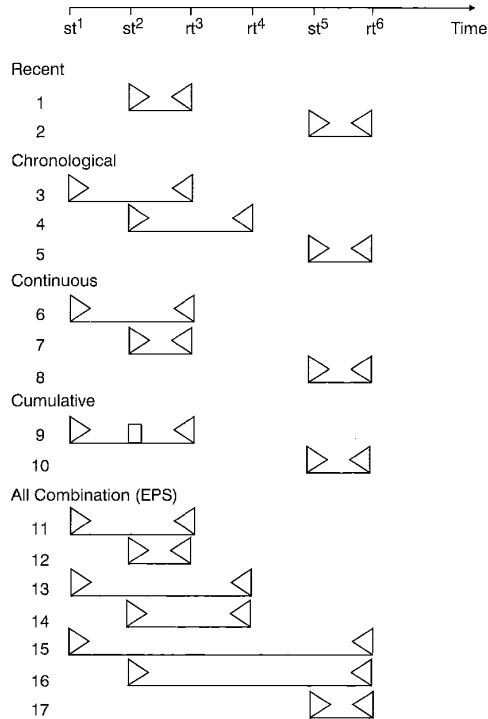
Considering the following order of instance arrivals (where j in e^j denotes the moment in which instance e has arrived):

$$[st, 1, 2]^1, [st, 2, 1]^2, [rt, 3, 1]^3, [rt, 2, 2]^4, [st, 3, 1]^5, [rt, 2, 1]^6,$$

it can be seen that the condition can only be detected by considering all combinations of events.

Fig. 6 shows which composite events would be generated in each of Snoop's modes (*recent*, *chronological*, *continuous* and *cumulative*) and EPS's generic mode (*all combination*). After processing the six primitive event instances in the generic mode the set of derived composite instances would be $(st1, rt3); (st2, rt3); (st1, rt4); (st2, rt4); (st1, rt6); (st2, rt6); (st5, rt6)$, (lines 11 through 17 in Fig. 6).

Now consider Snoop's modes where the upward propagation of events is determined by the instances that start (\triangleright) and finish (\triangleleft) the detection of a composite event (the box symbol (line 9) just represents the membership of a primitive instance in a composite instance). In *recent* mode the set of composite instance is $(st2, rt3); (st5, rt6)$. In *chronological* mode the set is $(st1, rt3); (st2, rt4); (st5, rt6)$. In the *continuous* mode the set is $(st1, rt3); (st2, rt3); (st5, rt6)$. Finally the *cumulative* mode the set is $(st2, st2, rt3); (st5, rt6)$.

**Fig. 6** Derivation of composite instances in Snoop and EPS

The union of the sets generated by each of Snoop's mode would have four elements plus the additional three-instance element (st1,st2,rt3). Thus the three other possible composite instances detected by EPS (lines 13, 15, and 16) would never be generated by Snoop, implying that such communication cycle would never be detected by Snoop.

As a consequence of adopting four specific modes, Snoop's algorithm just goes up in the tree. During the upward propagation of instances in the tree, the Snoop monitor removes instances from the lower levels according to the processing mode in use. This means that Snoop combines the EPS's first and the fourth phases in only one, while the second and the third phases are not necessary because a copy of each instance of the lower level is passed to the upper level. However, this approach is not suitable for generating all combinations (generic mode), and combinations can be lost.

9 Performance of EPS

Unfortunately, a mathematically rigorous analysis of the complexity of EPS's algorithm is very difficult because of the large number of factors and the difficulty of quantifying them. Based on the experience gained by testing EPS for various forms of shared event trees and different numbers of instances and arrival rates, we identified the main factors that determine the algorithm's performance. Below is a list in decreasing order of influence:

Types of operators: Among the possible types of operators, some contribute with more event instances in the tree than others. For example, the operator "||" does not cause the creation of many instances because it does not employ the cartesian product of the instances at the corresponding child nodes. However, if the majority of nodes in the tree is of type operator "&", operator ";" or repetition (operator **TimesDf** and operator **TimesEq**), the number of instances grows exponentially with the height of the tree.

Frequency of new primitive event instances: the more new event instances arrive at the EPS during a given period of time, the lower will be the overall performance of the algorithm. This is due to the fact that more instances remain in the tree and may be used when applying the cartesian product for deriving higher-level composite instances.

Size of the detection window: it determines the period of time during which an instance is used to generate new intermediate instances in the tree. Thus, if the detection window is large, more instances remain in the tree and hence cause the creation of even more instances.

Distribution of instances in the tree: this is also a factor that has influence on the total number of instances in the tree. If the distribution of instances at the leaf nodes is unbalanced, most intermediate instances will be concentrated on certain parts of the tree, and thus will contribute less to the generation of more intermediate instances. Thus an uniform distribution of instance arrivals for all event types is worse for the overall performance than a nonuniform distribution.

Structure of the tree: the tree structure also has a twofold effect on the generation of intermediate instances, i.e. on the total number of instances in a tree. For example, comparing a binary balanced tree and a backbone-like tree (a linear tree where each node has a single additional node) with a same distribution of instances at the leaf nodes, the first will have less intermediate nodes than the

last and hence will require less processing according to our algorithm.

Condition expression: the event condition associated with a R/N determines the number of instances that can be eliminated from the tree once the corresponding application node has been reached. For example, the more restrictive the condition is, the smaller is the chance of being able to satisfy the condition with a set of composite instances. Hence less instances are consumed or cancelled and remain in the tree.

Most of the factors depend on the nature of the applications using the EPS, such as, common structures and conditions associated with composite events of interest, primitive instance arrival rates for the different types of events, etc. However, fine tuning of parameters, such as the size of the detection window, or the periodicity of execution of the fourth phase, helps to improve the services performance. For example, if outdated instances are removed more frequently less instances are kept in the tree, which in turn reduces the amount of composite instances that are generated and processed.

10 Conclusions

Detection of composite events has been an active area of research in the database community, but more recently it is also being applied to other domains such as distributed computing, security and surveillance systems or analysis of financial markets.

Our interest in this subject arose from the area of distributed system monitoring and control, where we are developing an infrastructure for availability management of distributed programs. In particular, we required a service capable of detecting composite events, expressed in terms of primitive events detected by *sensors*, and describing complex relations of events concerning availability problems of software and hardware components. This lead to the design of a generic event detection service that can be used by any application program, e.g. a monitoring program, and which handles delays caused by network latency in the delivery of event messages from sensors.

The most significant contribution of this work is the use of shared trees to represent event expressions with common subexpressions. The major advantage of this optimisation is that it reduces significantly the cost of scanning the leaf nodes and evaluating subevents that are common to many event trees. With our technique the arrival of an instance of an event type appearing in a shared tree triggers only once the detection procedure for composite events.

As part of the design process we have formally described this event processing method using a customised attribute grammar. In this description we specified each node of the event tree as an object and the interactions related to each of the four processing phases as invocations of the corresponding methods at adjacent tree nodes. With this formal description the actual programming of the event processing algorithm was carried out without any major problems.

The event processing service described was fully implemented in Java with the use of the RMI facility as the basic means for communication among the EPS and its clients, i.e. the application processes and the sensors. (The EPS is public domain software and may be downloaded from <http://www.ime.usp.br/~moreto/eps/eps.tar>.) We chose the Java language because it is object-oriented, portable

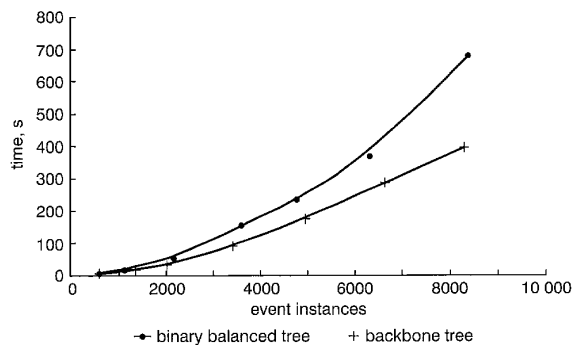


Fig. 7 Processing time of balanced binary tree against backbone-like tree

across many environments, and because RMI provides a simple method invocation interface for remote applications. Inheritance was very helpful for implementing the classes for the tree nodes (representing the primitive event types and operators) and for structuring the EPS code in general. We used threads to implement concurrency within the EPS, and this was also well supported by Java. Since many threads manipulate the event and instance trees, we were forced to synchronise the access to them using Java's synchronisation facility.

We tested the performance of the EPS for composite event expressions corresponding to extreme tree configurations. Fig. 7 shows the total processing time for a binary balanced tree versus a backbone-like tree for different numbers of stored instances, i.e. instances that had previously arrived at the EPS. We measured the processing time from the moment of arrival of the last instance necessary for a notification until the moment of the actual notification to the application process. This last-arriving instance was always corresponding to a deeply nested event type in the event expression. Unlike our initial belief the processing time with the backbone-like tree lead to a linear increase of time while the balanced tree resulted in an exponential behaviour. This occurred because although we kept the total number of instances the same, their distribution was different in each tree. And this fact apparently had a bigger influence on the total processing time than the shape of the tree.

We also compared the total processing time for a tree shared by two applications with the total processing time for two copies of the same tree, and for different distributions of previous instances in the trees. Again, here we measured the time between the arrival of the last instance and the moment of the second notification to the application program. Fig. 8 shows that the speedup gained when sharing trees is less for a binary balanced tree (tests 1–4) than for a backbone-like tree (test 5).

In all tests the processing of events using shared trees was faster than using simple trees, as expected. But this difference corresponds only to the time saved in the tree processing. In addition, it is important to consider also the time spent to find the trees which have to be processed. Then the total speedup gained is more than the one showed in the Fig. 8. Because related work does not present any data of the time spent to find the trees, this information could not be included in the graphics.

Our performance tests, and in particular Fig. 7, also showed that EPS is able to process approximately up to 1000–2000 event instances per minute, and therefore is suitable for systems with a moderate instance creation rate (by all sensors).




Fig. 8 Comparison of shared tree against duplicated tree

Although the performance of the EPS was very good, its implementation could be optimised and the results could be even better, meaning a larger number of instances processed by minute.

10.1 Future work

We are aware of the major limitations of our service, namely its centralised implementation and its assumptions about approximately synchronised machine clocks, and we plan to address these issues in future work.

A centralised implementation of any basic service embodies the problem of contention at concurrent access. To reduce the contention the best solution is to have several servers, each being in charge of monitoring events of a certain class. Although this reduces the contention at each server, it makes the detection of composite events much more complicated. Distributed event detection is an interesting and challenging problem which we plan to investigate in future.

One direction of future work is also to implement a class which acts as an interface between the EPS and the application-specific sensors. In addition to facilitating the programming of these sensors this class would also implement a clock synchronisation protocol among the communicating machines, making the EPS portable and less dependent on the existence of network clock synchronisation services.

Another line of future work is to investigate methods for identifying event subexpressions within a R/N that are already present in a previous event tree and implement simple transformations on event expressions to take advantage of the structure of an already existing event tree.

11 References

- 1 TSAI, J.J.P., and YANG, S.J.H.: 'Monitoring and debugging distributed real-time systems' (IEEE Computer Society Press, 1995)
- 2 TOMBROS, D., DITTRICH, K.R., GEPPERT, A., and KRADOLFE, M.: 'EvE, an event engine for distributed workflow enactment'. Technical report, Department of computer science, University of Zürich, 1996
- 3 MANSOURI-SAMANI, M.: 'Monitoring of distributed systems'. PhD thesis, Imperial College of Science, Technology and Medicine, University of London, 1995
- 4 ENDLER, M., and D'SOUZA, A.: 'Supporting distributed application management in SAMPa'. Proceedings of 3rd international conference on *Configurable distributed systems*, Annapolis, USA, IEEE Computer Society, May 1996, pp. 177–184
- 5 GARCIA, A.F., BEDER, D.M., and RUBIRA, C.M.E.: 'An exception handling mechanism for developing dependable object-oriented software based on meta-level approach'. Proceedings of ISSRE'99, Boca Raton, FL, IEEE, November 1999
- 6 GEPPERT, A., DITTRICH, K.R., and GATZIU, S.: 'The SAMOS active DBMS prototype'. Technical report 94.16, Institut für Informatik, University of Zürich, 1994

- 7 HSEUSH, W., and KAISER, G.: 'Modeling concurrency in parallel debugging', *ACM SIGPLAN Not.*, 1990, **25**, (3), pp. 11–20
- 8 JAGADISH, H.V., SHMUELII, O., and GEHANI, N.H.: 'COMPOSE: a system for composite event specification and detection'. Technical report, AT&T Bell Laboratories, 1992
- 9 CHAKRAVARTHY, S., and MISHRA, D.: 'Snoop: an expressive event specification language for active databases'. Technical report UF-CIS-TR-93-007, University of Florida, 1993
- 10 RAMAKRISHNA, Y.S., MELLIAH-SMITH, P.M., MOSER, L.E., DILLON, L.K., and KUTTY, G.: 'Interval logics and their decision procedures, part I: An interval logic', *Theor. Comput. Sci.*, 1996, **166**, (1 and 2), pp. 1–47